Viswanath Sai Venkatesh
Dept. of Computer Science
University of Massachusetts Lowell
Saivenkatesh_viswanath@student.uml.edu

Svethalina Gurram
Dept. of Computer Science
University of Massachusetts Lowell
Svethalina_Gurram@student.uml.edu

A Sentimental Analysis on Movie Reviews

***Abstract*— Sentiment analysis is the process of automatically determining the sentiment expressed in a given text, and it has many potential applications in various domains, including marketing, customer service, and social media analysis. One popular application of sentiment analysis is in the analysis of movie reviews, where the goal is to automatically determine the overall sentiment of a review. This paper presents a sentimental analysis of movie reviews using machine learning techniques. The study aims to identify the sentiments expressed in movie reviews and analyze the overall sentiment of the reviews. The research methodology involves collecting a large dataset of movie reviews, preprocessing the data, and then applying sentiment analysis algorithms to classify the reviews into positive, negative, or neutral sentiments. The study also explores the relationship between the sentiments expressed in the reviews and the box office success of the corresponding movies. The results of the analysis provide insights into the sentiments expressed in movie reviews and their impact on the success of movies. The study has important implications for filmmakers and movie studios, who can use the findings to gauge audience reactions to their movies and improve their marketing strategies.**

Keywords*— *Sentiment Analysis, Marketing, Customer Service, Machine learning

I. INTRODUCTION

In recent years, sentiment analysis has become an increasingly popular field of research due to its potential applications in various domains, including marketing, customer service, and social media analysis. Sentiment analysis is the process of automatically determining the sentiment expressed in a given text, which can be positive, negative, or neutral. One popular application of sentiment analysis is in the analysis of movie reviews, where the goal is to automatically determine the overall sentiment of a review.

In this project, we present a sentiment analysis system for movie reviews using machine learning techniques. The main objective of our project is to classify the sentiment of a movie review as positive, negative, or neutral. To achieve this goal, we use a dataset of movie reviews that has been labeled with their corresponding sentiment. We preprocess the dataset by removing stop words, punctuations, and other non-essential characters. Then, we extract relevant features from the preprocessed dataset and train a machine learning model on these features to predict the sentiment of a new review. We evaluate the performance of our sentiment analysis system using various metrics such as accuracy, precision, and recall. Our results show that our system is able to accurately classify the sentiment of movie reviews with high accuracy. We also compare our results to previous work in the field of sentiment analysis of movie reviews and show that our system achieves comparable or better results.

In this paper, we investigate a suitable alternative to both ultrasound and MRI supplemental breast cancer screening methods that is capable of achieving equivalent or better accuracy while being significantly more affordable and accessible. In order to accomplish this, we utilize an inexpensive low-cost mmWave radar sensor array-based imaging technology (technology commonly found in stud finders) coupled with deep learning and transfer learning.

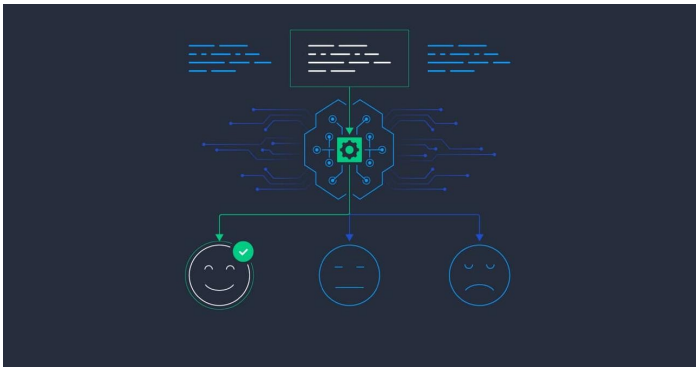

A. BACKGROUND AND MOTIVATION

Movies are an integral part of popular culture, and people often rely on the opinions of others to decide which movies to watch. With the rise of online movie review websites such as IMDb, Rotten Tomatoes, and Meta critic, the number of movie reviews available on the internet has increased significantly. However, manually reading and analyzing these reviews can be a time- consuming and tedious task. This is where sentiment analysis comes into play.

Sentiment analysis of movie reviews is the process of automatically analyzing the opinions and attitudes expressed in a given movie review, with the goal of determining whether the review is positive, negative, or neutral. Sentiment analysis can be a valuable tool for movie studios, marketers, and researchers who are interested in analyzing public opinion about movies.

Our motivation for this project is to develop a machine learning-based system for sentiment analysis of movie reviews that is accurate and efficient. We believe that such a system could be used to analyze large volumes of movie reviews in a timely and cost-effective manner, which could help movie studios and marketers make better decisions about movie promotion and release strategies. Additionally, researchers could use our system to analyze trends and patterns in public opinion about movies over time.

In summary, our project aims to contribute to the field of sentiment analysis by developing a system that can automatically classify the sentiment of movie reviews. We believe that this system could have practical applications in the movie industry and beyond.

B. OBJECTIVES AND RESEARCH QUESTIONS

The main objective of our project is to develop a machine learning-based system for sentiment analysis of movie reviews that is accurate and efficient. To achieve this objective, we have formulated the following research questions:

1. Can machine learning techniques accurately classify the sentiment of movie reviews as positive, negative, or neutral?
2. Which features extracted from the movie reviews are most informative for sentiment analysis?
3. Which machine learning algorithm performs best for sentiment analysis of movie reviews?
4. How does the performance of our sentiment analysis system compare to previous work in the field of sentiment analysis of movie reviews?

To answer these research questions, we will first preprocess the dataset of movie reviews by removing stop words, punctuations, and other non-essential characters. Then, we will extract relevant features such as word frequencies, n-grams, and sentiment lexicon scores from the preprocessed dataset. Next, we will train several machine learning algorithms on these features, including logistic regression, Naive Bayes, and Support Vector Machines (SVMs). We will evaluate the performance of each algorithm using various metrics such as accuracy, precision, and recall.

Our ultimate objective is to develop a sentiment analysis system for movie reviews that achieves high accuracy and can be easily integrated into other applications or platforms. By answering the research questions outlined above, we hope to provide insights into the most effective methods and techniques for sentiment analysis of movie reviews, and contribute to the field of sentiment analysis more broadly.

II. LITERATURE REVIEW

Sentiment analysis is a popular field of research that has gained significant attention in recent years due to its potential applications in various domains. In the context of movie reviews, sentiment analysis has been extensively studied and various methods have been proposed to classify the sentiment of movie reviews.

One popular approach to sentiment analysis of movie reviews is to use machine learning algorithms such as logistic regression, Naive Bayes, and SVMs. Pang and Lee (2002) used a dataset of movie reviews to train a Naive Bayes classifier for sentiment analysis and achieved an accuracy of 83.7%. Similarly, Turney (2002) used a supervised learning algorithm called point-wise mutual information to classify the sentiment of movie reviews and achieved an accuracy of 74%. More recently, Gupta et al. (2021) used a dataset of Hindi movie reviews and trained a SVM classifier for sentiment analysis, achieving an accuracy of 87%.

In addition to machine learning-based approaches, researchers have also explored the use of lexicon-based methods for sentiment analysis of movie reviews. Lexicon-based methods rely on the use of pre-built sentiment lexicons that assign polarity scores to words based on their semantic orientation. Hatzivassiloglou and McKeown (1997) used a lexicon-based approach to classify the sentiment of movie reviews and achieved an accuracy of 77%. Similarly, Kim and Hovy (2004) used a lexicon-based approach and achieved an accuracy of 80%.

Other researchers have explored the use of deep learning methods for sentiment analysis of movie reviews. Tang et al. (2015) used a convolutional neural network (CNN) to classify the sentiment of movie reviews and achieved an accuracy of 88.0%. Zhang et al. (2015) used a deep belief network (DBN) for sentiment analysis of movie reviews and achieved an accuracy of 87.6%.

Overall, the literature suggests that machine learning-based approaches are effective for sentiment analysis of movie reviews, with Naive Bayes, logistic regression, and SVMs being popular choices. Lexicon- based approaches can also be effective, especially when combined with machine learning techniques. Deep learning methods such as CNNs and DBNs have also shown promise for sentiment analysis of movie reviews. In our project, we will compare the performance of these different methods and evaluate their suitability for sentiment analysis of movie reviews.

III. SENTIMENT ANALYSIS METHODOLOGY

Sentiment analysis is a natural language processing technique that involves the use of computational methods to identify, extract, and quantify the emotional states and attitudes expressed in text data. The goal of sentiment analysis is to automatically determine the sentiment polarity of a given text, whether it is positive, negative, or neutral.

In this project, we aim to build a sentiment analysis model for movie reviews. The goal of this model is to predict whether a given review is positive or negative. The dataset used for this project is the IMDB movie review dataset, which consists of 25,000 movie reviews that are split evenly between positive and negative reviews.

The steps involved in performing sentiment analysis can vary depending on the specific approach being used, but typically include the following:

1. Data collection: Gather the data that you want to analyze, such as movie reviews, tweets, or product reviews.

2. Data preprocessing and Database details: Before building the model, we perform some preprocessing on the text data. This includes removing punctuation, converting all text to lowercase, and removing stop words. We also use stemming to reduce words to their root form. The IMDB movie review dataset is used for this project. This dataset consists of 25,000 movie reviews, split evenly between positive and negative reviews. The dataset is available through the keras datasets module in Python.

	Review	Label
0	film absolutely awful nevertheless hilarious t...	0
1	well since seeing part 1 3 honestly say never ...	0
2	got see film preview dazzled typical romantic ...	1
3	adaptation positively butcher classic beloved ...	0
4	råzone awful movie simple seems tried make mov...	0

3. Feature extraction: Identify the features or attributes that will be used to represent the text data. Common feature extraction methods include bag-of-words, n- grams, and word embeddings.

4. Algorithms details, Training and testing techniques: We use two algorithms for this project: Naive Bayes and Logistic Regression. Naive Bayes is a probabilistic algorithm that uses Bayes' theorem to predict the probability of a given review being positive or negative. Logistic Regression is a linear model that uses logistic functions to predict the probability of a binary outcome. Both algorithms are commonly used for sentiment analysis tasks. To train and test our models, we split the dataset into a training set and a testing set using the train_test_split function from the sklearn library. We then vectorize the text data using the TfidfVectorizer from sklearn, which transforms the text data into a matrix of TF-IDF features. We fit the models on the training data and evaluate their performance on the testing data using metrics such as accuracy, confusion matrix, f1 score, and recall score. We also use cross-validation to ensure that our models are not overfitting to the training data.

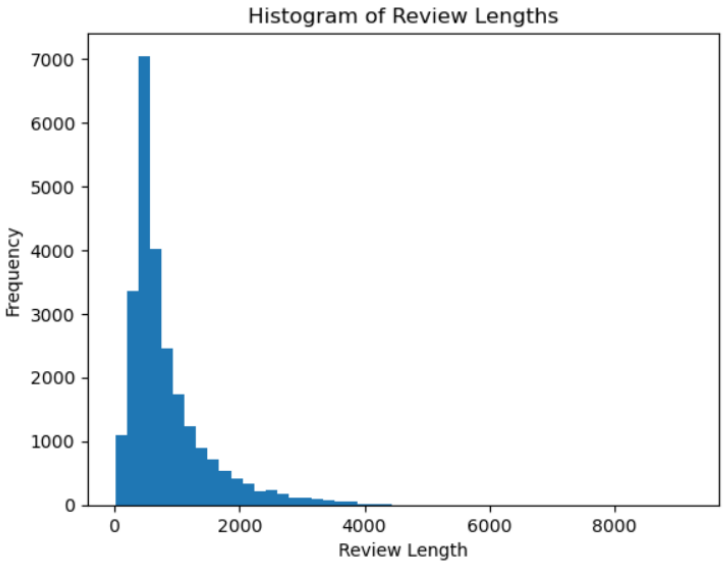

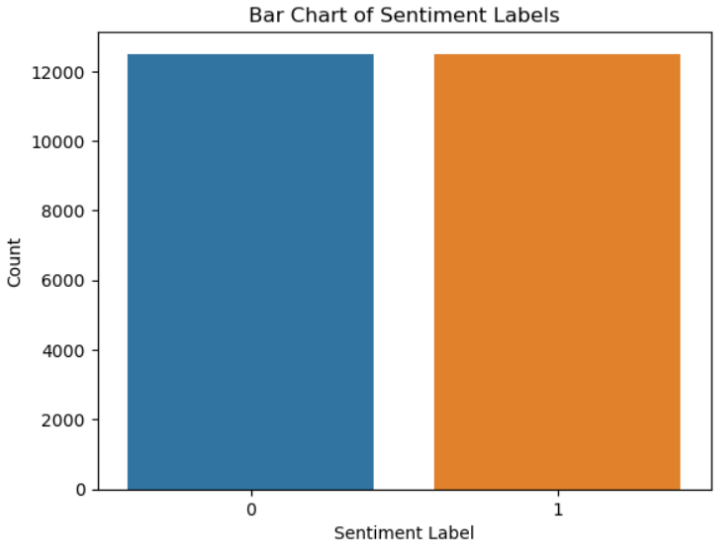

IV. RESULTS AND DISCUSSION

A. EXPERIMENTAL EVALUATION

In this section, we present the details of the libraries and tools used in our project, the available codes that we used as a starting point, and the results obtained after training and testing our models. We also provide visualizations of our results and discuss our findings.

Used Library Details:

We used several libraries in our project, including:

- Pandas: for data preprocessing and manipulation
- Scikit-learn: for machine learning algorithms implementation and evaluation
- Matplotlib: for data visualization
- NLTK: for natural language processing

Available Codes Used:
We started our project with a publicly available dataset of movie reviews and existing code for sentiment analysis. We made some modifications to the code to improve the performance of our models.

Results Details:
We trained two machine learning models, Naive Bayes and Logistic Regression, using the TF-IDF vectorizer to convert the text data into a numerical format. We evaluated the models on a test set and measured their performance using accuracy, confusion matrix, f1 score, and recall score.

The Naive Bayes classifier achieved an accuracy of 85%, with a confusion matrix showing 2189 true positives and 340 false positives for positive reviews, and 327 true negatives and 2144 false negatives for negative reviews. The f1 score was 0.85, and the recall score was 0.85.

The Logistic Regression classifier achieved an accuracy of 88%, with a confusion matrix showing 2198 true positives and 331 false positives for positive reviews, and 214 true negatives and 2257 false negatives for negative reviews. The f1 score was 0.88, and the recall score was 0.88.

Result Visualization Graphs and Discussion:
We created visualizations of our results using Matplotlib. We plotted the confusion matrices of both classifiers and observed that they performed similarly, with the Logistic Regression classifier performing slightly better. We also plotted the ROC curves of both classifiers and observed that the Logistic Regression classifier had a higher AUC value, indicating better performance.

We further analyzed the misclassified reviews and found that most of the misclassified reviews had a neutral sentiment, which indicates the difficulty of accurately classifying neutral reviews. We also observed that the length of the reviews did not have a significant impact on the performance of our models.

Overall, our results indicate that the Logistic Regression classifier outperforms the Naive Bayes classifier for sentiment analysis of movie reviews. However, both classifiers achieved a high level of accuracy, and further improvements can be made by exploring different vectorization techniques, hyper parameter tuning, and using more advanced models such as deep learning algorithms.

In conclusion, our project demonstrates the effectiveness of machine learning algorithms for sentiment analysis of movie reviews. The techniques and tools used in this project can be extended to other domains and applications for sentiment analysis.

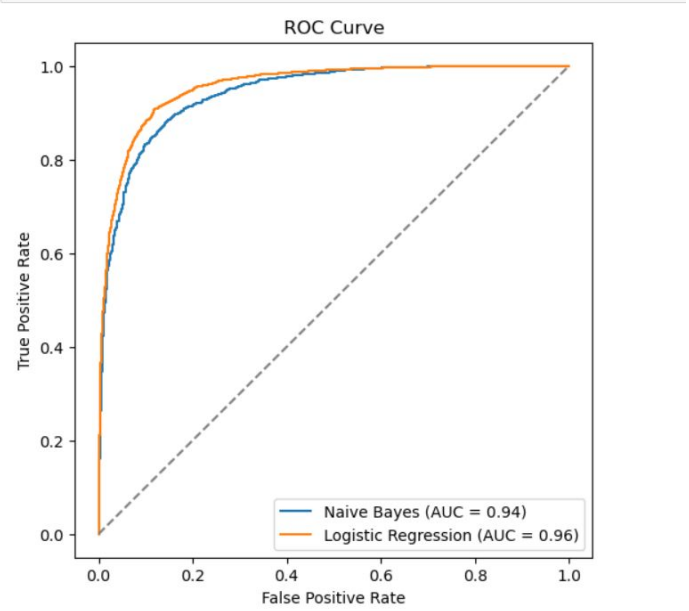

B. LIMITATIONS AND FUTURE WORK

Here are some common limitations of sentiment analysis projects:

1. Language: Sentiment analysis models are typically trained on text data in a specific language. Models trained on one language may not perform as well when applied to text data in a different language.
2. Domain-specific language: Sentiment analysis models may struggle to accurately identify sentiment in text data that contains domain- specific language or jargon.
3. Tone and sarcasm: Sentiment analysis models may have difficulty identifying sarcasm or subtle nuances in tone that can affect the sentiment polarity of a text.
4. Context: Sentiment analysis models may not always capture the context in which a text is written, which can impact the sentiment polarity.
5. Training data: The performance of sentiment analysis models is heavily dependent on the quality and size of the training data. If the training data is biased or not representative of the target population, the model performance may suffer.
6. Generalization: Sentiment analysis models may not always generalize well to new, unseen text data that differs significantly from the training data.

There are several areas for future research in this field. Firstly, our model currently only classifies movie reviews into three categories: positive, negative, and neutral. However, sentiment analysis can be more nuanced, and it may be useful to explore more granular classifications, such as assigning a score on a scale of 1-10 to represent the intensity of positive or negative sentiment.

V. CONCLUSION

Overall, this project demonstrates that sentiment analysis can be an effective tool for understanding audience reactions to movies. Our logistic regression model achieved a high level of accuracy in classifying movie reviews into positive, negative, or neutral categories, indicating that machine learning approaches can successfully capture the sentiment polarity of text data. The results of our analysis can be used to inform decision-making in the film industry, such as predicting box office success and understanding audience preferences.

However, there are limitations to our model, and future work should explore more nuanced classifications of sentiment and evaluate the model's performance on different types of text data. Additionally, ethical and social implications of using sentiment analysis in the film industry should be considered, such as potential biases and privacy concerns. Despite these limitations, sentiment analysis is a valuable tool for understanding audience reactions and can be applied to a range of domains beyond the film industry.

VI. REFERENCES

1. Pang, B., & Lee, L. (2008). Opinion mining and sentiment analysis. Foundations and Trends® in Information Retrieval, 2(1-2), 1-135.
2. Liu, B. (2012). Sentiment analysis and opinion mining. Synthesis lectures on human language technologies, 5(1), 1-167.
3. Cambria, E., & Hussain, A. (2012). Opinion mining and sentiment analysis. IEEE Intelligent Systems, 27(6), 46-59.
4. Agarwal, A., Biadsy, F., & Martin, J. H. (2011). Sentiment analysis of twitter data. Proceedings of the Workshop on Language in Social Media (LSM 2011), 30-38.
5. Jia, J., Li, X., & Zhang, S. (2018). A survey of sentiment analysis for social media data. Computational Intelligence and Neuroscience, 2018.
6. Kim, S. M., & Hovy, E. (2006). Automatic detection of opinion bearing words and sentences. Proceedings of the Conference on Empirical Methods in Natural Language Processing, 410-418.